


GenAI Delivery Excellence

Skills Validation Workshop




Your Instructor


- Background
- Position
- Organization




Etiquette



Mute microphone




No recording




Ask questions

Course format



Instructor Lectures
(10%)



Challenge Labs
(90%)


Why should I participate in this Skills Validation Workshop?

- 01 Validate the knowledge you gained from the Google generative AI training program
- 02 Prove your ability to implement generative AI solutions for customers
- 03 Demonstrate a broad understanding of using Google Cloud tools to implement real-world generative AI use cases
- 04 Get more work for you and your company!



Target audience

This workshop assumes intermediate to advanced knowledge of generative AI, Google Cloud, Google Gemini, and related APIs



Google Cloud

- ✓ ML Engineers
- ✓ Developers
- ✓ Cloud Architects

Prerequisites

- ✓ Generative AI Foundational training in Google Cloud Skills Boost
- ✓ Experience deploying solutions to Google Cloud
- ✓ Working knowledge of Google BigQuery, Cloud Storage, Agent Platform, and related APIs
- ✓ Programming experience with Python and Jupyter Notebooks





Agenda



Day 1

- 9:00 Introduction
- 9:30 Challenge 1: Gemini Prompt Security
- 11:30 Challenge 2: Programming a RAG System in BigQuery
- 2:00 Challenge 3: Testing and Evaluation
- 4:00 Bonus Challenge 4: Using Agent Builder

Day 2

- 9:00 Challenge 5: Alaska Department of Snow Online Agent

Grading

Challenges	Points
Challenge One: Gemini Prompt Security	20
Challenge Two: Programming a RAG System in BigQuery	20
Challenge Three: Testing and Evaluation	20
Bonus Challenge Four: Using Agent Builder	10
Challenge Five: Alaska Department of Snow Online Agent	40
Total points possible:	110

- 80 points is required to pass the workshop

Submitting your work ...

01 Log into **GitHub** (if you do not have an account, create one)

02 Create a public GitHub repository named:
genai-skills-workshop-<<your-name>>

03 Most of the challenges will be done using Jupyter Notebooks.
You will need to upload those notebooks to this repository for grading.

04 Make this repo available to your instructor.



Open Google Cloud Skills Boost

01 Open an incognito window (or private/anonymous window).

02 Go to the Qwiklabs URL cloudskillsboost.google

03 Sign in with existing account, or join with new account, using the **email you used to register** for the course.

04 Launch the course from **In Progress** section of your home page.

Jumpstart your cloud career

Not sure where to start? Find featured learning below. We give you temporary credentials to actual cloud resources, so you can learn the cloud using the real thing.

In Progress



SCHEDULED COURSE

Welcome to Cloud Digital Leader

Dec 7, 2021 9:00AM EST

Virtual

Access issues

If you encounter any problems while setting up your Qwiklabs account, follow the troubleshooting steps here: <http://traininghelp.cloudlearning.io>



Challenge One

Gemini Prompt Security

System instructions example

System instructions



You are a programming chat bot.

Answering coding questions using Python unless another language is specified.

Format code according to the Python PEP 8 Style guide.

Explain the job to
the model

USER Write an Apache Beam Pipeline that reads a CSV file with the fields: ID, Customer, Address, Credit_Limit, creates a Hash of the ID field, and writes the records to a BigQuery table with the same fields.

AI

```
import apache_beam as beam
from apache_beam.options.pipeline_options import PipelineOptions
import hashlib

# Replace with your actual BigQuery project ID and dataset ID
PROJECT_ID = 'your-project-id'
DATASET_ID = 'your_dataset_id'
TABLE_ID = 'your_table_id'

def hash_id(element):
    """Hashes the ID field using SHA-256."""
    element['ID'] = hashlib.sha256(str(element['ID']).encode('utf-8')).hexdigest()
    return element

def run():
```

It is good practice to give systems instructions on what *not* to do

System instructions ?  


×

You are a programming chat bot.
Answering coding questions using Python unless another language is specified.
Format code according to the Python PEP 8 Style guide.
Do Not answer questions unrelated to Computers or Programming.
For any question unrelated to Computers or Programming, just say "I can't help with that."

USER Write me a 1500 word essay on the life of William Shakespeare


AI I can't help with that.

Prompt injection attempts to get the model to ignore its instructions



The latest models won't be tricked by simple prompt injection tactics

When building an application that uses LLMs, you can add guardrails and filters between the user and the model



Challenge one: Gemini prompt security

Goal: Demonstrate how to program a secure and safe generative AI system

Requirements:

- ✓ Using Python, create a chat application that uses the latest version of Gemini.
- ✓ Create system instructions that include chatbot goals and restrictions.
- ✓ Implement prompt filtering to validate user input.
 - Recommended solution is to use the Google Model Armor API for prompt injection and jailbreak detection
- ✓ Implement Gemini safety filters.
 - **Bonus:** Use the Google Model Armor API and the Sensitive Data Protection API for response filtering
- ✓ Validate model responses for safety, and only return responses that are deemed ok.

Challenge one: Instructions

- 1 Log into your Cloud Skills Boost environment provided by your instructor.
- 2 Create a Jupyter Notebook using Agent Platform Colab Enterprise.
- 3 Create a coding and IT chatbot with appropriate system instructions and restrictions.
- 4 Implement the requirements as outlined on the previous slide.
- 5 Make sure to use text blocks and comments to document your solution.
- 6 Upload the Jupyter Notebook to your GitHub repo and send the link to the instructor for grading.

Tips



- In the Google Cloud console, go to **Agent Platform | Create Prompt**. Use the console to experiment with your prompt and click the **Get code** link for sample code.
- Model Armor documentation:
 - [Model Armor Overview](#)
 - [Sanitize prompts and responses](#)

A decorative graphic on the left side of the slide. It includes a green circle at the top left, a red rounded rectangle with a wavy pattern at the top right, a large yellow rounded rectangle with a wavy pattern in the center containing the text '02', a blue rounded rectangle at the bottom left, and a green rounded rectangle with a wavy pattern at the bottom right.


02

Challenge Two

Programming a RAG System in BigQuery

Retrieval Augmented Generation (RAG)

- RAG allows the system to search a repository of information for data relevant to a user question and pass that data to the LLM with the question



In BigQuery ML, you can connect to an external AI model

```
CREATE OR REPLACE MODEL `CustomerReview.Embeddings`  
  REMOTE WITH CONNECTION `us.embedding_conn`  
  OPTIONS (ENDPOINT = 'text-embedding-005');
```

Use the ML.GENERATE_EMBEDDING function to create embedding vectors using SQL

```
CREATE OR REPLACE TABLE
`CustomerReview.customer_reviews_embedded` AS
SELECT *
FROM ML.GENERATE_EMBEDDING(
    MODEL `CustomerReview.Embeddings`,
    (SELECT review_text AS content FROM
    `CustomerReview.customer_reviews`)
);
```

Use the VECTOR_SEARCH function to do a nearest neighbor search

```
SELECT query.query,base.content
FROM
  VECTOR_SEARCH(
    TABLE `CustomerReview.customer_reviews_embedded`,
    'ml_generate_embedding_result',
    (
      SELECT ml_generate_embedding_result,content AS query
      FROM
        ML.GENERATE_EMBEDDING(MODEL `CustomerReview.Embeddings`,
          (SELECT 'service' AS content))),
    top_k => 5,
    options => '{"fraction_lists_to_search": 0.01}');
```


Challenge two: Programming a RAG system in BigQuery

Goal: Demonstrate your ability to program a RAG system that uses BigQuery to generate embeddings and perform a vector search.

Requirements:

- ✓ Students are provided with a set of frequently asked questions for the fictitious town of Aurora Bay, Alaska
- ✓ They will need to load the data into a BigQuery table.
- ✓ Students need to create embeddings for each record and store them in BigQuery.
- ✓ Implement a chatbot that can search the BigQuery embeddings to answer user questions accurately using the proprietary data.


Challenge two: Instructions

- 1 Create a Jupyter Notebook using Agent Platform Colab Enterprise.
- 2 Import the following file into a BigQuery table:
`gs://labs.roitraining.com/aurora-bay-faqs/aurora-bay-faqs.csv`
- 3 Generate embeddings for each question-and-answer pair, store the embeddings along with the existing fields.
- 4 Program a chatbot that can use a vector search to find the data required to answer the user's question.
- 5 Pass the data and question to Gemini to generate an accurate response.
- 6 Make sure this is coded in Python and well-documented in your Notebook.
- 7 Upload the Notebook to GitHub for grading.

Tips



- Here is a Cloud Skills Boost lab that does a similar task: https://www.cloudskillsboost.google/catalog_lab/31904
- BigQuery generative AI documentation: <https://cloud.google.com/bigquery/docs/generative-ai-overview>



03

Challenge Three

Testing and Evaluation

Functions that return results from LLMs need to be tested like any other function

- Functions that return deterministic results are simple
 - Using an LLM for classification or sentiment analysis are examples
 - Tested like any other function
 - The response from the model is either right or wrong
- Model responses that are indeterminate are more difficult to test
 - There is no single right answer to: “Write me a response to this email”
 - To automate testing use an LLM to evaluate correctness

LLM classification example

```
def getPositiveOrNegative(prompt):  
    response = model.generate_content(  
        """Context: You look at messages and categorize them as  
        positive, negative, or neutral.
```

Output only Positive, Negative, or Neutral.

Message: {0}.

```
Output: """.format(prompt)  
)
```

```
return response.text.strip()
```

```
print(getPositiveOrNegative("Dinner was cold and undercooked."))
```



Negative

Unit testing a deterministic function

```
import unittest

class TestPositiveOrNegative(unittest.TestCase):
    def test_getPositiveOrNegative1(self):
        response = getPositiveOrNegative("Dinner was Great")
        self.assertEqual(response, "Positive")

    def test_getPositiveOrNegative2(self):
        response = getPositiveOrNegative("That broccoli was undercooked and cold")
        self.assertEqual(response, "Negative")

    def test_getPositiveOrNegative3(self):
        response = getPositiveOrNegative("We went to the new italian place for dinner")
        self.assertEqual(response, "Neutral")
```

```
unittest.main(argv=[''], verbosity=2, exit=False)
```

```
test_getPositiveOrNegative1 (__main__.TestPositiveOrNegative) ... ok
test_getPositiveOrNegative2 (__main__.TestPositiveOrNegative) ... ok
test_getPositiveOrNegative3 (__main__.TestPositiveOrNegative) ... ok
```

Ran 3 tests in 0.768s

OK

LLM text generation example

```
def writeTweet(prompt):  
    response = model.generate_content(  
        """Context: You write Tweets for the Marketing Department at Luigi's Italian Cafe.  
  
        1. Keep your Tweets below 100 characters  
        2. Include the hashtag #EatAtLuigis at the end of every tweet  
  
        Input: Write me a for our All You Can eat Spaghetti Wednesdays deal.  
        Output: Spaghetti Wednesday is here! Come get all you can eat spaghetti  
               and meatballs for only $9.99! #EatAtLuigis  
  
        Input: {0}  
        Output: """.format(prompt)  
    )  
  
    return response.text.strip()
```


Ask the LLM if the generated content followed the rules

```
def does_tweet_follow_rules(tweet):  
    response = model.generate_content(  
        """Does the tweet follow the following rules:  
  
        1. Keep your Tweets below 100 characters  
        2. Include the hashtag #EatAtLuigis at the end of every tweet  
  
        Only return Yes or No  
  
        Tweet: {0}  
        Output: """.format(tweet)  
    )  
  
    return response.text.strip()
```

Run unit tests to see if the generated tweet followed the rules the LLM was given

```
import unittest

class TestTweetRules(unittest.TestCase):

    def test_tweet_results_1(self):
        generated_tweet = writeTweet("Write a tweet about our Thanksgiving Day Special")
        correct = does_tweet_follow_rules(generated_tweet)
        self.assertEqual(correct, "Yes")

    def test_does_not_follow_rules(self):
        generated_tweet = "Thanksgiving feast at Luigi's! Enjoy our special menu."
        correct = does_tweet_follow_rules(generated_tweet)
        self.assertEqual(correct, "No")
```

Google GenAI evaluation service lets you evaluate your large language models

01

Pointwise Evaluations

Evaluate a single model based on one or more metrics

02

Pairwise Evaluations

Evaluate two models (or prompts) based on one or more metrics to select the preferred one.

GenAI evaluation service supports two types of metrics

01

Computed metrics

Mathematical methods that compare model responses to a source of truth

02

Model-based metrics

Model-based metrics that use rules and a large language model to determine the quality of responses


Computed metrics

Metric	Description	Use cases
Exact Match	- Does the prediction (response from the model) match the reference source exactly?	Useful when responses are constrained to exact, structured answers
Rouge	- Recall-Oriented Understudy for Gisting Evaluation - Evaluates the overlap between n-grams, word sequences, or word pairs between the generated text and reference text	Summarization tasks, as it captures both the recall and the relevance of generated text.
Bleu	- Bilingual Evaluation Understudy - Compares the overlap of n-grams between the generated text and reference text - Higher BLEU scores indicate closer similarity to reference texts	Translation or summarization where n-gram similarity is relevant

Model-based metrics include...

Metric	Description	Output
Fluency	Assess a response's language mastery	Integer between 1 and 5 (5 being the best)
Coherence	Assess a response's ability to provide a coherent, easy-to-follow reply.	Integer between 1 and 5 (5 being the best)
Safety	Assess a response's level of safety	0 is unsafe, 1 if safe
Groundedness	Assess a response's ability to provide or reference information included only in the input text	0 is ungrounded, 1 if grounded
Fulfillment	Assess a response's ability to completely fulfill instructions	Integer between 1 and 5 (5 being the best)
Summarization Quality	Assess a response's overall ability to summarize text	Integer between 1 and 5 (5 being the best)
Custom Metric	You provide the rules that the model should use to evaluate response quality	

Using the Evaluation API



Creating an evaluation dataset

```
contexts = [str(record) for record in apartment_records]
full_prompts = [prompt + str(record) for record in apartment_records]

eval_dataset = pd.DataFrame(
    {
        "content": full_prompts,
        "instruction": full_prompts,
        "context": contexts,
    }
)
```


Creating an evaluation task

```
qa_eval_task = EvalTask(  
    dataset=eval_dataset,  
    metrics=["fulfillment", "groundedness"],  
    experiment="apartment-listing-generation",  
)
```

Run the evaluation

```
run_ts = datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
result = qa_eval_task.evaluate(
    model=model,
    experiment_run_name=f"apartment-listing-gen-{run_ts}"
)

evaluation_results = []
evaluation_results.append(result)
```

Examine the results

```
display_eval_report(("Eval Result", result.summary_metrics, result.metrics_table))
```

Eval Result

Summary Metrics

	row_count	fulfillment/mean	fulfillment/std	groundedness/mean	groundedness/std
0	10.0	5.0	0.0	1.0	0.0

Report Metrics

	content	instruction	context	response	fulfillment/explanation	fulfillment/confidence	fulfillment	groundedness/e
0	Write a one-paragraph apartment listing\nto pr...	Write a one-paragraph apartment listing\nto pr...	{'Address': '123 West 14th Street, New York, N...	Luxury living awaits in this spacious 2-bedroo...	STEP 1: Assess instruction understanding: The ...	1.0	5.0	STEP 1: accurately r
1	Write a one-paragraph apartment listing\nto pr...	Write a one-paragraph apartment listing\nto pr...	{'Address': '456 East 57th Street, New York, N...	Live comfortably in this charming 789 sq ft 1-...	STEP 1: Assess instruction understanding: The ...	1.0	5.0	STEP 1: All a respons

Challenge three: Testing and evaluation

Goal: Demonstrate your ability to write tests and evaluate responses from large language models.

Requirements:

- ✓ Create functions that perform specific tasks utilizing Google Gemini
- ✓ Write unit tests to ensure the functions perform as expected
- ✓ Use the Google Evaluation API to evaluate and compare functions using various different prompt.

Challenge three: Instructions

- 1 Create a Jupyter Notebook using Agent Platform Colab Enterprise.
- 2 Create a Python function that uses Gemini to classify user questions into one of the following categories: **Employment**, **General Information**, **Emergency Services**, or **Tax Related**
- 3 Create a second function that generates social media posts for government announcements like weather emergencies, holidays, school closings, etc.
- 4 Write unit tests for each function using pytest.
- 5 Use the Google Evaluation API to evaluate and compare Gemini responses from different prompts.
- 6 Upload the Notebook to GitHub for grading.

Tips



- Here is a Cloud Skills Boost lab that demonstrates testing GenAI functions
https://partner.cloudskillsboost.google/catalog_lab/31628
- Here is a Cloud Skills Boost lab that uses the Evaluation API
https://partner.cloudskillsboost.google/catalog_lab/31631
- Gen AI Evaluation Service documentation:
<https://cloud.google.com/vertex-ai/generative-ai/docs/model-reference/evaluation>

A decorative graphic on the left side of the slide. It features a green shape at the top left with a black line pattern, a yellow star-like shape at the top right, a red diamond shape in the center with a white line pattern and the number '04' in white, a yellow circle at the bottom left, and a blue rounded rectangle at the bottom right with a black line pattern.

04

Bonus Challenge Four

Building Agents with AI Applications

Instructor demo: Using AI Applications

Conversation agents



Conversational agent PREVIEW

Build a conversational agent using both rules and generative AI that responds naturally, and predictably.

[CREATE](#)

Goal

Goal*

You are a Chatbot for the Alsaka Town of Aurora Bay. Your name is Alice.
You answer questions from citizens, tourists, and anyone else who wants information about the town and its services.
Do not answer questions unrelated to the town. If users ask questions completely unrelated to the town, say you don't know.

Instructions

[Templates](#)

Instructions

- 1 - Greet the user. Tell them who you are, what you can do, and ask what they would like.
- 2 - When answering questions you can search the data store: `${TOOL:Aurora Bay Data Store}`
- 3 - Do not just make things up. If you cannot find the answer to a question, say you don't know.

What kind of data are you importing?

For more information, see [Prepare data for ingesting](#).

- ☒ Unstructured documents (PDF, HTML, TXT and more)
Supported file formats: PDF, HTML, TXT (CSV for FAQ, DOCX)
- ☐ Linked unstructured documents (JSONL with metadata)
JSONL files with unstructured document links and its metadata [requirements](#)
- ☐ Structured FAQ data for a chat application (CSV)
Structured FAQ data

Select a folder or a file you want to import

FOLDER

FILE

gs:// *

[labs.roitraining.com/aurora-bay-faqs](#)

Bonus challenge four: Building agents with AI Applications

Goal: Demonstrate your ability to use Google AI Applications to create a Conversational Agent that that uses a Playbook and Data Store to answer user questions.

Requirements:

- ✓ Create a conversational agent using AI Applications.
- ✓ Create an AI Applications data store.
- ✓ Create a playbook with a goal, instructions, and data store that answers user questions.

Instructions

In this challenge, you will create an agent for the fictional town of Aurora Bay, Alaska. The agent's job will be to answer frequently asked questions from citizens.

- 1 In the Google Cloud console, go to **AI Applications**.
- 2 Create a new **Conversational Agent** app called **Aurora Bay Agent** in US-Central1.
- 3 In the Default Playbook, add an appropriate goal and instructions.
- 4 Add a data store that is connected to the following Google Cloud Storage bucket:
gs://labs.roitraining.com/aurora-bay-faqs
- 5 Test the agent. Make sure it answers questions based on the data from the data store and does not just make things up.
- 6 Export the agent as JSON and upload it to your GitHub repository for grading.



05

Challenge Five

Alaska Department of Snow Online Agent






The Alaska Department of Snow (ADS) serves 750,000 people across 650,000 square miles and relies on interagency communication to deliver services. During snow forecasts, regional offices face high call volumes with questions about plowing, school closures, and other disruptions. ADS is exploring an online agent or chatbot to offload routine inquiries, but some administrators have reservations about cloud services, and the CFO is concerned about costs.

Existing tech environment

- ADS's technical infrastructure is a mix of old and new solutions, some of which are fully vendor-managed, and some of which are managed by the ADS team.
- Most information is stored online in web pages, text files, or PDFs.
- Use of generative AI is not sanctioned by the department at this point, but many individual staff began using GenAI tools to support their own individual productivity, before the statewide IT department temporarily paused use earlier this year.
- The department is open to using generative AI solutions if they can be proven accurate, reliable, safe, and cost-effective.



Case study requirements

- You have met with the team at ADS proposing the use of an online agent (chatbot) to alleviate call volume and address common resident questions.
- You need to build a prototype chatbot, and demonstrate functionality to the leadership team and decision-makers.
- In your presentation, you should also be prepared to:
 - Review the architecture
 - Address questions and objections related to security, privacy, and accuracy
 - Demonstrate a working proof of concept



Challenge five: Alaska Department of Snow Online Agent

Goal: Demonstrate your ability to create a secure, accurate, production-quality generative AI agent that can be deployed online

Requirements:

- ✓ Backend data store for RAG
- ✓ Access to backend API functionality
- ✓ Unit tests for agent functionality
- ✓ Evaluation data using the Google Evaluation service API
- ✓ Prompt filtering and response validation implemented into the solution
- ✓ Log all prompts and responses
- ✓ Generative AI agent deployed to a website

Instructions

- 1 Create a diagram that depicts your solution to the Alaska Snow Department case study
- 2 Implement the case study any way you like as long as it meets the requirements outlined in Challenge Five
- 3 Create the backend RAG system using data from the following Cloud Storage bucket:
`gs://labs.roitraining.com/alaska-dept-of-snow`
- 4 Don't forget to include tests and evaluation data
- 5 Upload all artifacts to your GitHub repository for grading



Thank you.